A Handbook of Statistical Analyses Using **R** — 2nd Edition

Brian S. Everitt and Torsten Hothorn

Density Estimation: Erupting Geysers and Star Clusters

8.1 Introduction

8.2 Density Estimation

The three kernel functions are implemented in R as shown in lines 1–3 of Figure 8.1. For some grid x, the kernel functions are plotted using the R statements in lines 5–11 (Figure 8.1).

The kernel estimator $\hat{f}$ is a sum of 'bumps' placed at the observations. The kernel function determines the shape of the bumps while the window width h determines their width. Figure 8.2 (redrawn from a similar plot in Silverman, 1986) shows the individual bumps $n^{-1}h^{-1}K((x-x_i)/h)$, as well as the estimate $\hat{f}$ obtained by adding them up for an artificial set of data points

```
R> x <- c(0, 1, 1.1, 1.5, 1.9, 2.8, 2.9, 3.5)
R> n <- length(x)
```

For a grid

```
R> xgrid <- seq(from = min(x) - 1, to = max(x) + 1, by = 0.01)
```

on the real line, we can compute the contribution of each measurement in x, with $h = 0.4$, by the Gaussian kernel (defined in Figure 8.1, line 3) as follows;

```
R> h <- 0.4
R> bumps <- sapply(x, function(a) gauss((xgrid - a)/h)/(n * h))
```

A plot of the individual bumps and their sum, the kernel density estimate $\hat{f}$, is shown in Figure 8.2.

8.3 Analysis Using R

8.3.1 A Parametric Density Estimate for the Old Faithful Data

```
R> logL <- function(param, x) {
+     d1 <- dnorm(x, mean = param[2], sd = param[3])
+     d2 <- dnorm(x, mean = param[4], sd = param[5])
+     -sum(log(param[1] * d1 + (1 - param[1]) * d2))
+ }
R> startparam <- c(p = 0.5, mu1 = 50, sd1 = 3, mu2 = 80, sd2 = 3)
R> opp <- optim(startparam, logL, x = faithful$waiting,
```

```
1  R> rec <- function(x) (abs(x) < 1) * 0.5
2  R> tri <- function(x) (abs(x) < 1) * (1 - abs(x))
3  R> gauss <- function(x) 1/sqrt(2*pi) * exp(-(x^2)/2)
4  R> x <- seq(from = -3, to = 3, by = 0.001)
5  R> plot(x, rec(x), type = "l", ylim = c(0,1), lty = 1,
6  +        ylab = expression(K(x)))
7  R> lines(x, tri(x), lty = 2)
8  R> lines(x, gauss(x), lty = 3)
9  R> legend(-3, 0.8, legend = c("Rectangular", "Triangular",
10 +          "Gaussian"), lty = 1:3, title = "kernel functions",
11 +          bty = "n")
```
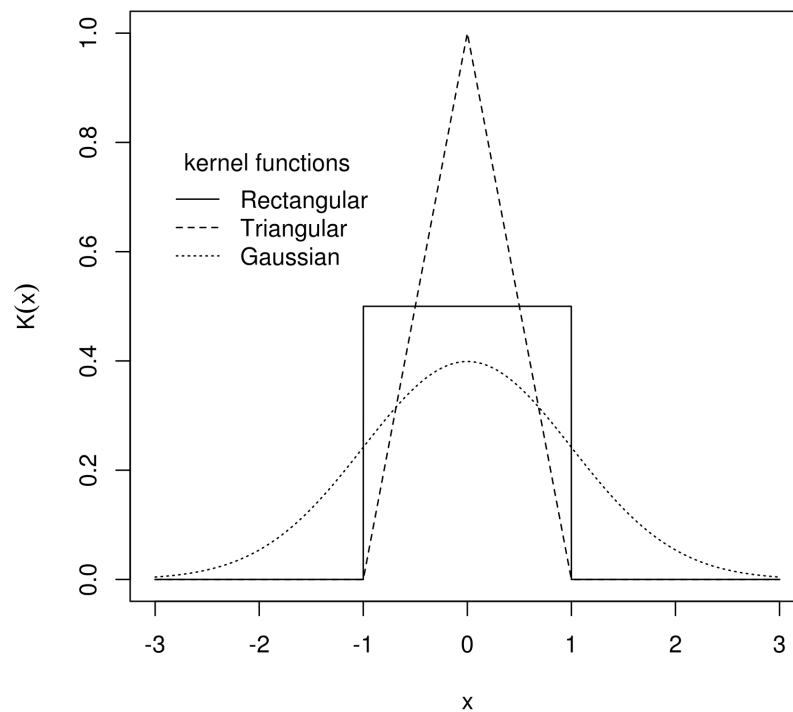

Figure 8.1   Three commonly used kernel functions.

```
1   R> plot(xgrid, rowSums(bumps), ylab = expression(hat(f)(x)),
2   +        type = "l", xlab = "x", lwd = 2)
3   R> rug(x, lwd = 2)
4   R> out <- apply(bumps, 2, function(b) lines(xgrid, b))
```
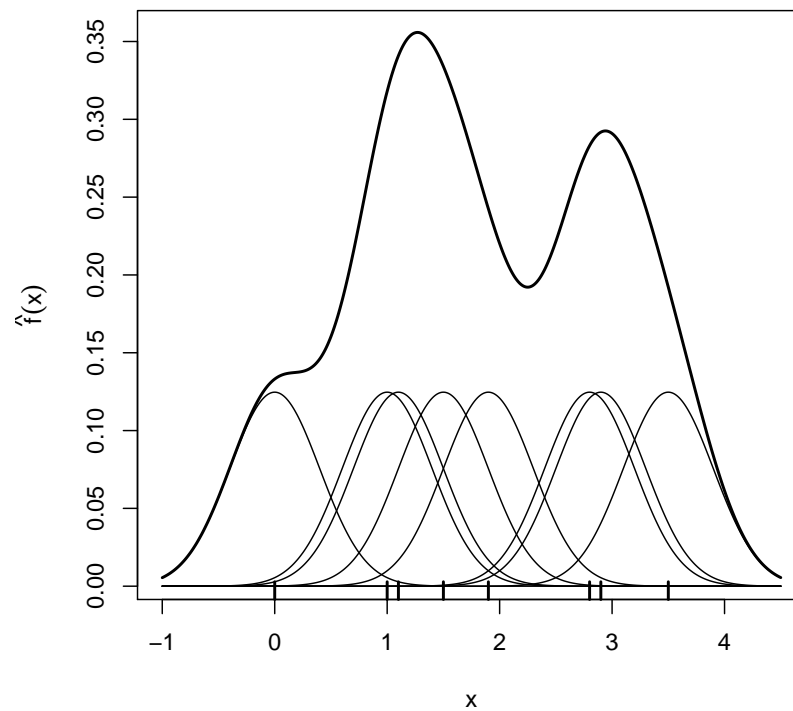

Figure 8.2   Kernel estimate showing the contributions of Gaussian kernels evaluated for the individual observations with bandwidth $h = 0.4$.

```
+                 method = "L-BFGS-B",
+                 lower = c(0.01, rep(1, 4)),
+                 upper = c(0.99, rep(200, 4)))
R> opp

$par
     p     mu1     sd1     mu2     sd2
 0.361  54.612   5.872  80.093   5.867

$value
[1] 1034
```

```
R> epa <- function(x, y)
+       ((x^2 + y^2) < 1) * 2/pi * (1 - x^2 - y^2)
R> x <- seq(from = -1.1, to = 1.1, by = 0.05)
R> epavals <- sapply(x, function(a) epa(a, x))
R> persp(x = x, y = x, z = epavals, xlab = "x", ylab = "y",
+         zlab = expression(K(x, y)), theta = -35, axes = TRUE,
+         box = TRUE)
```
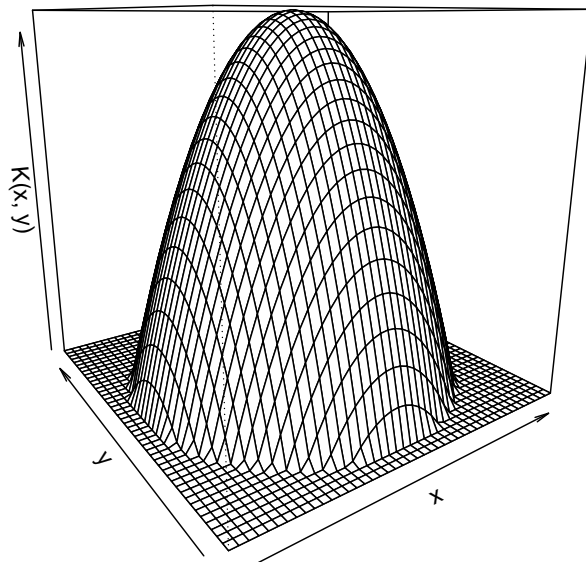

Figure 8.3   Epanechnikov kernel for a grid between $(-1.1, -1.1)$ and $(1.1, 1.1)$.

```
1  R> data("faithful", package = "datasets")
2  R> x <- faithful$waiting
3  R> layout(matrix(1:3, ncol = 3))
4  R> hist(x, xlab = "Waiting times (in min.)", ylab = "Frequency",
5  +        probability = TRUE, main = "Gaussian kernel",
6  +        border = "gray")
7  R> lines(density(x, width = 12), lwd = 2)
8  R> rug(x)
9  R> hist(x, xlab = "Waiting times (in min.)", ylab = "Frequency",
10 +        probability = TRUE, main = "Rectangular kernel",
11 +        border = "gray")
12 R> lines(density(x, width = 12, window = "rectangular"), lwd = 2)
13 R> rug(x)
14 R> hist(x, xlab = "Waiting times (in min.)", ylab = "Frequency",
15 +        probability = TRUE, main = "Triangular kernel",
16 +        border = "gray")
17 R> lines(density(x, width = 12, window = "triangular"), lwd = 2)
18 R> rug(x)
```
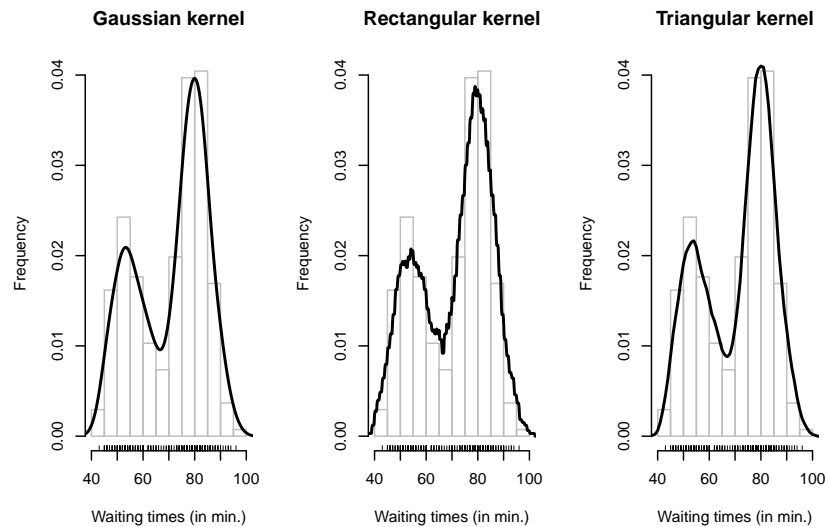

Figure 8.4   Density estimates of the geyser eruption data imposed on a histogram of the data.

```
R> library("KernSmooth")
R> data("CYGOB1", package = "HSAUR2")
R> CYGOB1d <- bkde2D(CYGOB1, bandwidth = sapply(CYGOB1, dpik))
R> contour(x = CYGOB1d$x1, y = CYGOB1d$x2, z = CYGOB1d$fhat,
+          xlab = "log surface temperature",
+          ylab = "log light intensity")
```
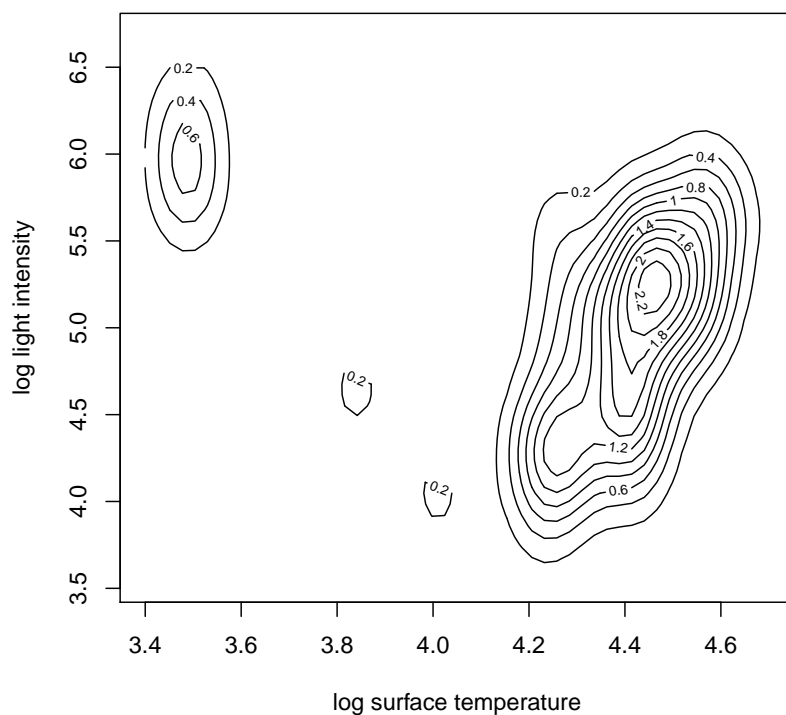

Figure 8.5   A contour plot of the bivariate density estimate of the CYGOB1 data,
i.e., a two-dimensional graphical display for a three-dimensional prob-
lem.

```
R> persp(x = CYGOB1d$x1, y = CYGOB1d$x2, z = CYGOB1d$fhat,
+          xlab = "log surface temperature",
+          ylab = "log light intensity",
+          zlab = "estimated density",
+          theta = -35, axes = TRUE, box = TRUE)
```
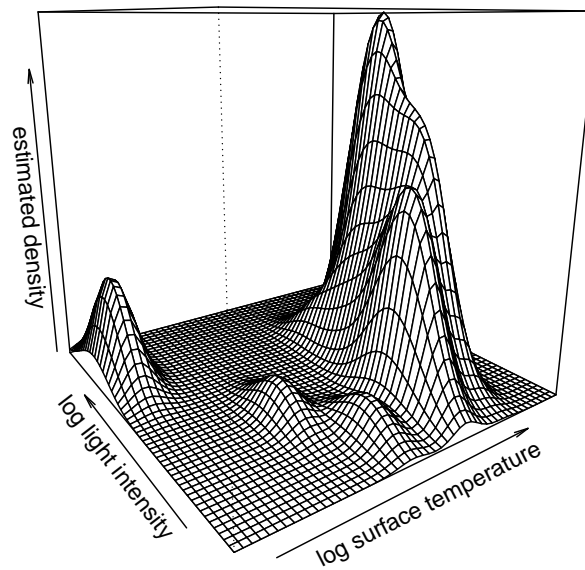

Figure 8.6   The bivariate density estimate of the CYGOB1 data, here shown in a three-dimensional fashion using the persp function.

```
$counts
function gradient
      55       55

$convergence
[1] 0
```

Of course, optimising the appropriate likelihood 'by hand' is not very convenient. In fact, (at least) two packages offer high-level functionality for esti-

mating mixture models. The first one is package **mclust** (Fraley et al., 2012) implementing the methodology described in Fraley and Raftery (2002). Here, a Bayesian information criterion (BIC) is applied to choose the form of the mixture model:

```
R> library("mclust")
R> mc <- Mclust(faithful$waiting)
R> mc
```

```
'Mclust' model object:
 best model: univariate, equal variance (E) with 2 components
```

and the estimated means are

```
R> mc$parameters$mean
```

```
   1    2
54.6 80.1
```

with estimated standard deviation (found to be equal within both groups)

```
R> sqrt(mc$parameters$variance$sigmasq)
```

```
[1] 5.87
```

The proportion is $\hat{p} = 0.36$. The second package is called **flexmix** whose functionality is described by Leisch (2004). A mixture of two normals can be fitted using

```
R> library("flexmix")
R> fl <- flexmix(waiting ~ 1, data = faithful, k = 2)
```

with $\hat{p} = 0.36$ and estimated parameters

```
R> parameters(fl, component = 1)
```

```
                 Comp.1
coef.(Intercept)   54.6
sigma               5.9
```

```
R> parameters(fl, component = 2)
```

```
                 Comp.2
coef.(Intercept)  80.10
sigma              5.87
```

We can get standard errors for the five parameter estimates by using a bootstrap approach (see Efron and Tibshirani, 1993). The original data are slightly perturbed by drawing n out of n observations *with replacement* and those artificial replications of the original data are called *bootstrap samples*. Now, we can fit the mixture for each bootstrap sample and assess the variability of the estimates, for example using confidence intervals. Some suitable R code based on the Mclust function follows. First, we define a function that, for a bootstrap sample indx, fits a two-component mixture model and returns $\hat{p}$ and the estimated means (note that we need to make sure that we always get an estimate of p, not $1 - p$):

```
R> opar <- as.list(opp$par)
R> rx <- seq(from = 40, to = 110, by = 0.1)
R> d1 <- dnorm(rx, mean = opar$mu1, sd = opar$sd1)
R> d2 <- dnorm(rx, mean = opar$mu2, sd = opar$sd2)
R> f <- opar$p * d1 + (1 - opar$p) * d2
R> hist(x, probability = TRUE, xlab = "Waiting times (in min.)",
+        border = "gray", xlim = range(rx), ylim = c(0, 0.06),
+        main = "")
R> lines(rx, f, lwd = 2)
R> lines(rx, dnorm(rx, mean = mean(x), sd = sd(x)), lty = 2,
+        lwd = 2)
R> legend(50, 0.06, lty = 1:2, bty = "n",
+         legend = c("Fitted two-component mixture density",
+                     "Fitted single normal density"))
```
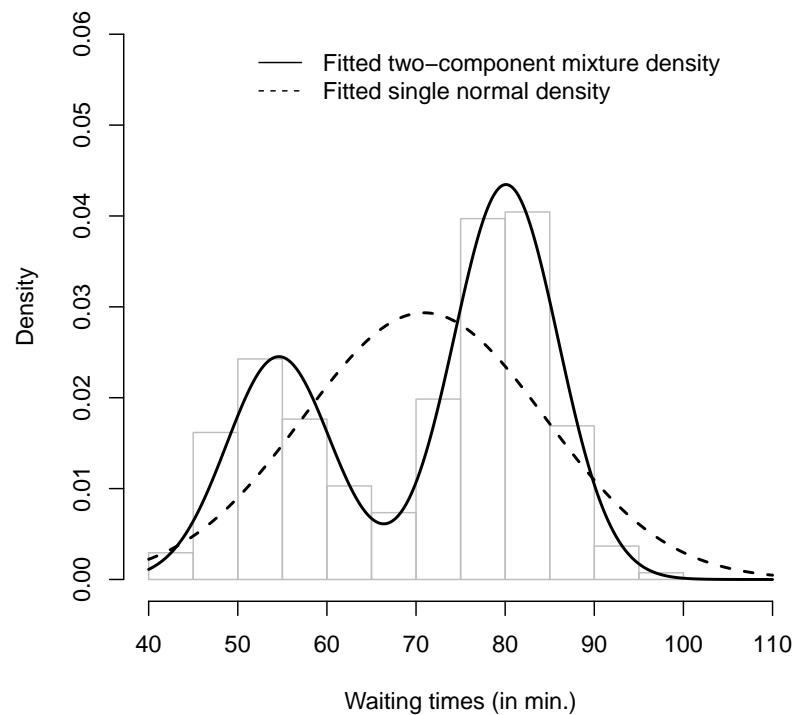

Figure 8.7 Fitted normal density and two-component normal mixture for geyser eruption data.

```
R> library("boot")
R> fit <- function(x, indx) {
+       a <- Mclust(x[indx], minG = 2, maxG = 2,
+                   modelNames="E")$parameters
+       if (a$pro[1] < 0.5)
+           return(c(p = a$pro[1], mu1 = a$mean[1],
+                                  mu2 = a$mean[2]))
+       return(c(p = 1 - a$pro[1], mu1 = a$mean[2],
+                                  mu2 = a$mean[1]))
+   }
```

The function `fit` can now be fed into the `boot` function (Canty and Ripley, 2012) for bootstrapping (here 1000 bootstrap samples are drawn)

```
R> bootpara <- boot(faithful$waiting, fit, R = 1000)
```

We assess the variability of our estimates $\hat{p}$ by means of adjusted bootstrap percentile (BCa) confidence intervals, which for $\hat{p}$ can be obtained from

```
R> boot.ci(bootpara, type = "bca", index = 1)
```

```
BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
Based on 1000 bootstrap replicates

CALL :
boot.ci(boot.out = bootpara, type = "bca", index = 1)

Intervals :
Level       BCa
95%    ( 0.304,  0.423 )
Calculations and Intervals on Original Scale
```

We see that there is a reasonable variability in the mixture model; however, the means in the two components are rather stable, as can be seen from

```
R> boot.ci(bootpara, type = "bca", index = 2)
```

```
BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
Based on 1000 bootstrap replicates

CALL :
boot.ci(boot.out = bootpara, type = "bca", index = 2)

Intervals :
Level       BCa
95%    (53.4, 56.1 )
Calculations and Intervals on Original Scale
```

for $\hat{\mu}_1$ and for $\hat{\mu}_2$ from

```
R> boot.ci(bootpara, type = "bca", index = 3)
```

```
BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
Based on 1000 bootstrap replicates
```

```
CALL :
boot.ci(boot.out = bootpara, type = "bca", index = 3)

Intervals :
Level       BCa
95%   (79, 81 )
Calculations and Intervals on Original Scale
```

Finally, we show a graphical representation of both the bootstrap distribution of the mean estimates *and* the corresponding confidence intervals. For convenience, we define a function for plotting, namely

```
R> bootplot <- function(b, index, main = "") {
+       dens <- density(b$t[,index])
+       ci <- boot.ci(b, type = "bca", index = index)$bca[4:5]
+       est <- b$t0[index]
+       plot(dens, main = main)
+       y <- max(dens$y) / 10
+       segments(ci[1], y, ci[2], y, lty = 2)
+       points(ci[1], y, pch = "(")
+       points(ci[2], y, pch = ")")
+       points(est, y, pch = 19)
+   }
```

The element `t` of an object created by `boot` contains the bootstrap replications of our estimates, i.e., the values computed by `fit` for each of the 1000 bootstrap samples of the geyser data. First, we plot a simple density estimate and then construct a line representing the confidence interval. We apply this function to the bootstrap distributions of our estimates $\hat{\mu}_1$ and $\hat{\mu}_2$ in Figure 8.8.

```
R> layout(matrix(1:2, ncol = 2))
R> bootplot(bootpara, 2, main = expression(mu[1]))
R> bootplot(bootpara, 3, main = expression(mu[2]))
```
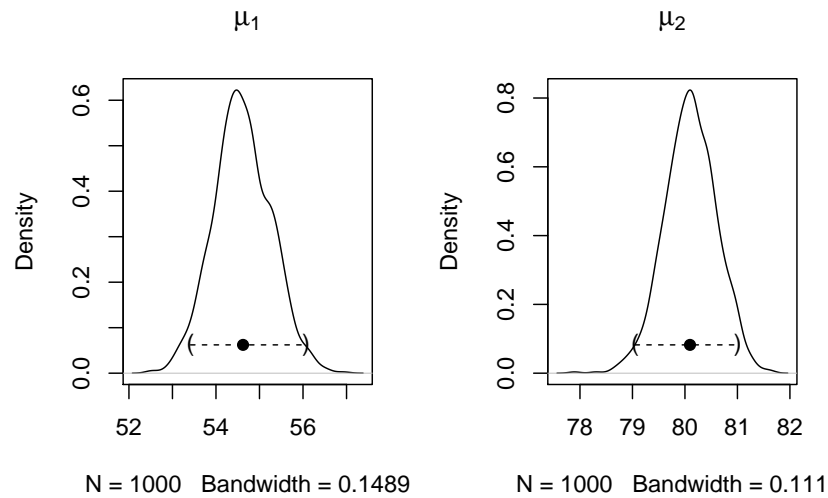

Figure 8.8 Bootstrap distribution and confidence intervals for the mean estimates of a two-component mixture for the geyser data.

Bibliography

Canty, A. and Ripley, B. D. (2012), **boot**: *Bootstrap R (S-PLUS) Functions*, URL `http://CRAN.R-project.org/package=boot`, R package version 1.3-7.

Efron, B. and Tibshirani, R. J. (1993), *An Introduction to the Bootstrap*, London, UK: Chapman & Hall/CRC.

Fraley, C. and Raftery, A. E. (2002), "Model-based clustering, discriminant analysis, and density estimation," *Journal of the American Statistical Association*, 97, 611–631.

Fraley, C., Raftery, A. E., and Wehrens, R. (2012), **mclust**: *Model-based Cluster Analysis*, URL `http://www.stat.washington.edu/mclust`, R package version 3.4.11.

Leisch, F. (2004), "FlexMix: A general framework for finite mixture models and latent class regression in R," *Journal of Statistical Software*, 11, URL `http://www.jstatsoft.org/v11/i08/`.

Silverman, B. (1986), *Density Estimation*, London, UK: Chapman & Hall/CRC.